



Part 21: Kubernetes Real-Time Troubleshooting

Introduction

Welcome to the world of Kubernetes troubleshooting, where every challenge is an opportunity to sharpen your skills and emerge victorious. Join us as we embark on a journey through common real-time scenarios, unraveling mysteries, and uncovering solutions along the way.



PART 21 - KUBERNETES REAL-TIME TROUBLESHOOTING

- Creating a Namespace and Resource-Limited Pod
- Managing Kubernetes Contexts
- Deploying a Pod on a Specific Node
- Scaling StatefulSets
- Implementing Readiness Probes

Scenario 101: Creating a Namespace and Resource-Limited Pod in Kubernetes

Scenario: You need to create a new namespace called limit and deploy a pod named resource-checker with specific CPU and memory requests and limits.

Solution:

1. **Create the Namespace:**
``kubectl create ns limit``
- **Define the Pod:**
apiVersion: v1
kind: Pod
metadata:
 namespace: limit
 labels:
 run: resource-checker
 name: resource-checker



```
spec:
  containers:
  - image: httpd:alpine
    name: my-container
  resources:
    requests:
      memory: "30Mi"
      cpu: "30m"
    limits:
      memory: "30Mi"
      cpu: "300m"
  dnsPolicy: ClusterFirst
  restartPolicy: Always
```

- **Deploy the Pod:**
``kubectl apply -f <pod-definition-file>.yaml``

Outcome: This setup ensures that the resource-checker pod has controlled resource usage, preventing it from consuming excessive CPU or memory.

Scenario 102: Managing Kubernetes Contexts and Current Context Retrieval

Scenario: You need to list all Kubernetes contexts and retrieve the current context using both kubectl commands and shell scripts.

Solution:

1. **List All Contexts:**
``kubectl config get-contexts``
- **Save Contexts to a File:**
``kubectl config get-contexts -o name > /root/filesystem/tmp``
- **Script to Get Current Context Using kubectl:**
``echo "kubectl config current-context" > current_context_using_kubectl.sh``
- **Script to Get Current Context Without kubectl:**
``echo "cat ~/.kube/config | grep current-context | sed 's/current-context: //' > current_context_without_kubectl.sh``

Outcome: These scripts provide flexibility in retrieving the current Kubernetes context, useful for automation and monitoring.

Scenario 103: Deploying a Pod on a Specific Node with Taints

Scenario: You need to deploy a pod on a specific node and ensure it remains scheduled there by managing taints.

Solution:

1. **Switch Context:**
``kubectl config use-context k8s-c1-H``
- **Create the Pod Definition:**
`apiVersion: v1`



```
kind: Pod
metadata:
  labels:
    run: pod1
    name: pod1
spec:
  nodeName: controlplane
  containers:
  - image: httpd:2.4.41-alpine
    name: pod1-container
  dnsPolicy: ClusterFirst
  restartPolicy: Always
```

- **Remove Taint from the Node:**

```
`kubectl taint nodes controlplane node-role.kubernetes.io/control-plane:NoSchedule-`
```

- **Deploy the Pod:**

```
`kubectl apply -f <pod-definition-file>.yaml`
```

Outcome: The pod is successfully deployed and scheduled on the specified node, ensuring consistent resource allocation.

Scenario 104: Scaling StatefulSets in Kubernetes

Scenario: You need to create and scale StatefulSets in a Kubernetes namespace to manage stateful applications efficiently.

Solution:

1. **Create the Namespace:**

```
`kubectl create ns project`
```

- **Define the StatefulSet:**

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-statefulset
  namespace: project
spec:
  serviceName: "my-service"
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
      - name: my-container
        image: nginx:latest
        ports:
```



- containerPort: 80

- **Scale Down the StatefulSet:**

```
`kubectl scale statefulset my-statefulset -n project --replicas=1`
```

Outcome: StatefulSets provide stable network identities and persistent storage, crucial for stateful applications.

Scenario 105: Implementing Readiness Probes with Service Dependencies

Scenario: You need to ensure a pod becomes ready only when a dependent service is available.

Solution:

1. **Create the First Pod with Readiness Probe:**

```
apiVersion: v1
kind: Pod
metadata:
  name: ready-if-service-ready
spec:
  containers:
  - name: nginx
    image: nginx:1.16.1-alpine
    readinessProbe:
      exec:
        command:
        - wget
        - -T2
        - -O-
        - http://service-am-i-ready:80
```

- **Create the Second Pod with Labels:**

```
apiVersion: v1
kind: Pod
metadata:
  name: am-i-ready
  labels:
    id: cross-server-ready
spec:
  containers:
  - name: nginx
    image: nginx:latest
```

- **Patch the Service to Select the Second Pod:**

```
`kubectl patch service service-am-i-ready --type=json' -p='[{"op": "add", "path": "/spec/selector", "value": {"id": "cross-server-ready"}}]`
```

Outcome: The first pod becomes ready only when the dependent service is available, ensuring proper service dependencies.



In the up-coming parts, we will discussion on more troubleshooting steps for the different Kubernetes based scenarios. So, stay tuned for the and follow @Prasad Suman Mohan for more such posts.

