



# Part 22: Kubernetes Real-Time Troubleshooting

## Introduction

Welcome to the world of Kubernetes troubleshooting, where every challenge is an opportunity to sharpen your skills and emerge victorious. Join us as we embark on a journey through common real-time scenarios, unraveling mysteries, and uncovering solutions along the way.

### Scenario 106: Managing Persistent Volumes and Claims in Kubernetes

**Scenario:** You need to create a PersistentVolume (PV) and PersistentVolumeClaim (PVC) to provide persistent storage to a deployment.

#### Solution:

##### 1. Create the PersistentVolume:

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: safari-pv
spec:
  capacity:
```

<https://www.linkedin.com/in/prasad-suman-mohan>



```
storage: 2Gi
accessModes:
  - ReadWriteOnce
persistentVolumeReclaimPolicy: Retain
hostPath:
  path: /Volumes/Data
```

## 2. Create the PersistentVolumeClaim:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: safari-pvc
  namespace: project-tiger
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
```

## 3. Create the Deployment Using the PVC:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: safari
  namespace: project-tiger
spec:
  replicas: 1
  template:
    spec:
      containers:
        - name: safari-container
          image: httpd:2.4.41-alpine
          volumeMounts:
            - name: safari-storage
              mountPath: /tmp/safari-data
      volumes:
        - name: safari-storage
          persistentVolumeClaim:
            claimName: safari-pvc
```

**Outcome:** The deployment successfully uses persistent storage, ensuring data persistence across pod restarts.

---

## Scenario 107: Installing and Configuring Metrics Server in Kubernetes

**Scenario:** You need to install the Metrics Server to monitor resource usage in your Kubernetes cluster.

**Solution:****1. Install Metrics Server:**

```
kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

**2. Allow Insecure TLS (If Required):**

```
kubectl patch deployment metrics-server -n kube-system --type=json -p=[{"op": "add", "path": "/spec/template/spec/containers/0/args/-", "value": "--kubelet-insecure-tls"}]
```

**3. Verify Metrics Server Installation:**

```
kubectl get deployment metrics-server -n kube-system  
kubectl top nodes
```

**Outcome:** The Metrics Server provides valuable insights into resource usage, aiding in cluster management and optimization.

---

**Scenario 108: Sorting Pods by Age and UID**

**Scenario:** You need to list pods sorted by their creation timestamp and UID for better management and monitoring.

**Solution:****Script to List Pods Sorted by Age:**

```
kubectl get pods --all-namespaces --sort-by=.metadata.creationTimestamp
```

**Script to List Pods Sorted by UID:**

```
kubectl get pods --all-namespaces --sort-by=.metadata.uid
```

**Outcome:** These scripts help in quickly identifying the oldest or newest pods, aiding in resource management and troubleshooting.

---

**Scenario 109: Managing Kube-Scheduler & scheduling pod**

**Scenario:** Use context: `kubectl config use-context k8s-c2-AC`. Ssh into the master node with `ssh cluster2-master1`. Temporarily stop the kube-scheduler, this means in a way that you can start it again afterwards. Create a single Pod named `manual-schedule` of image `httpd:2.4-alpine`, confirm its created but not scheduled on any node. Now you're the scheduler and have all its power, manually schedule that Pod on node `cluster2-master1`. Make sure it's running. Start the kube-scheduler again and confirm its running correctly by creating a second Pod named `manual-schedule2` of image `httpd:2.4-alpine` and check if it's running on `cluster2-worker1`.



## Solution:

### Context:

```
kubectl config use-context k8s-c2-AC
```

### Steps:

#### 1. SSH into the Master Node:

- Use ssh cluster2-master1 to access the master node.

#### 2. Temporarily Stop the Kube-Scheduler:

- Move the kube-scheduler.yaml manifest to temporarily stop the scheduler:

```
cd /etc/kubernetes/manifests/ && mv kube-scheduler.yaml ../
```

#### 3. Create a Pod Named manual-schedule:

- Run the following command to create the pod:

```
kubectl run manual-schedule --image=httpd:2.4-alpine
```

- Confirm the pod is created but not scheduled:

```
kubectl get pods
```

#### 4. Manually Schedule the Pod:

- Edit the pod's YAML to include nodeName: cluster2-master1 and apply the changes:

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: manual-schedule
```

```
spec:
```

```
  nodeName: cluster2-master1
```

```
  containers:
```

```
  - name: manual-schedule
```

```
    image: httpd:2.4-alpine
```

- Apply the manifest:

```
kubectl apply -f <edited-pod-yaml>
```

#### 5. Restart the Kube-Scheduler:

- Move the kube-scheduler.yaml back to its original location:

```
mv ../kube-scheduler.yaml .
```

- Verify the scheduler is running by creating another pod:

```
kubectl run manual-schedule2 --image=httpd:2.4-alpine
```

```
kubectl get pods -o wide
```

### Outcome:

- The manual-schedule pod should be running on cluster2-master1.



- The manual-schedule2 pod should be scheduled and running on cluster2-worker1.

## Scenario 110: Creating a ServiceAccount with Role and RoleBinding

**Scenario:** Use context: `kubectl config use-context k8s-c1-H` Create a new ServiceAccount processor in Namespace project-hamster. Create a Role and RoleBinding, both named processor as well. These should allow the new SA to only create Secrets and ConfigMaps in that Namespace.

### Solution:

Context:

```
kubectl config use-context k8s-c1-H
```

#### Steps:

##### 1. Create a ServiceAccount:

- Create a ServiceAccount named processor in the project-hamster namespace:

```
kubectl create serviceaccount processor -n project-hamster
```

##### 2. Create a Role:

- Define a Role named processor that allows creating Secrets and ConfigMaps:

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: Role
```

```
metadata:
```

```
  namespace: project-hamster
```

```
  name: processor
```

```
rules:
```

```
- apiGroups: [""]
```

```
  resources: ["secrets", "configmaps"]
```

```
  verbs: ["create"]
```

- Apply the Role definition:

```
kubectl apply -f processor-role.yaml
```

##### 3. Create a RoleBinding:

- Define a RoleBinding to bind the ServiceAccount to the Role:

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: RoleBinding
```

```
metadata:
```

```
  name: processor
```

```
  namespace: project-hamster
```

```
subjects:
```

```
- kind: ServiceAccount
```

```
  name: processor
```

```
  namespace: project-hamster
```

```
roleRef:
```



```
kind: Role
name: processor
apiGroup: rbac.authorization.k8s.io
  ○ Apply the RoleBinding definition:
```

```
kubectl apply -f processor-rolebinding.yaml
```

**Outcome:**

- The processor ServiceAccount in the project-hamster namespace will have permissions to create Secrets and ConfigMaps.

---

By following these steps, you can create a well-structured Word document that clearly outlines each scenario, its steps, and expected outcomes. If you need further assistance with formatting or specific Word features, let me know!



In the up-coming parts, we will discussion on more troubleshooting steps for the different Kubernetes based scenarios. So, stay tuned for the and follow @Prasad Suman Mohan for more such posts.

