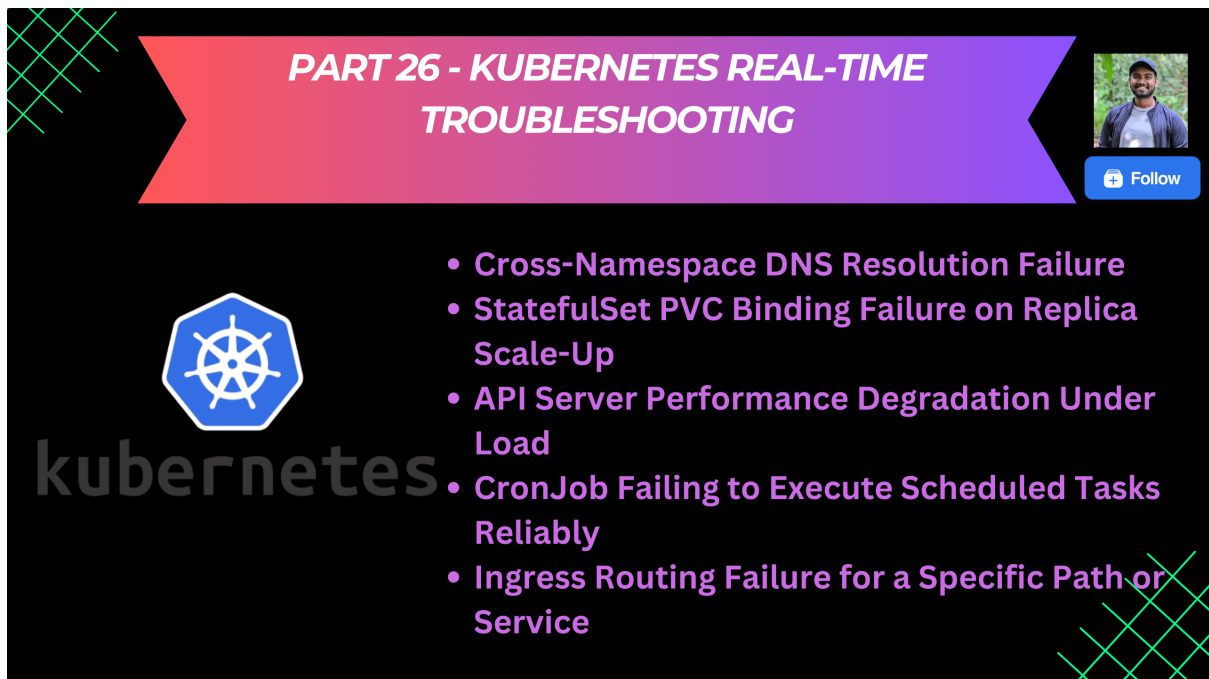




# Part 26: Kubernetes Real-Time Troubleshooting

## Introduction

Welcome to the world of Kubernetes troubleshooting, where every challenge is an opportunity to sharpen your skills and emerge victorious. Join us as we embark on a journey through common real-time scenarios, unraveling mysteries, and uncovering solutions along the way.



**PART 26 - KUBERNETES REAL-TIME TROUBLESHOOTING**

- Cross-Namespace DNS Resolution Failure
- StatefulSet PVC Binding Failure on Replica Scale-Up
- API Server Performance Degradation Under Load
- CronJob Failing to Execute Scheduled Tasks Reliably
- Ingress Routing Failure for a Specific Path or Service

### Scenario 126: Cross-Namespace DNS Resolution Failure

#### Scenario:

Pods running in the `project-alpha` namespace are suddenly unable to resolve Services hosted within the `project-beta` namespace (e.g., `beta-service.project-beta.svc.cluster.local`). Internal DNS resolution for services within `project-alpha` works correctly, and pods in `project-beta` have no issues. This is disrupting inter-service communication vital for a multi-component application.

#### Solution

<https://www.linkedin.com/in/prasad-suman-mohan>



## Context:

kubectl config use-context k8s-cl7-prod

## Steps:

### 1. Verify DNS Resolution Failure:

Exec into a pod in `project-alpha`:

```
kubectl exec -it -n project-alpha <alpha-pod-name> -- /bin/sh
```

Attempt DNS lookups using `nslookup` or `dig` (install if needed: `apk add bind-tools` or `apt-get update && apt-get install dnsutils`):

# Should fail or timeout

```
nslookup beta-service.project-beta.svc.cluster.local
```

# Should succeed

```
nslookup alpha-service.project-alpha.svc.cluster.local
```

# Should succeed

```
nslookup kubernetes.default.svc.cluster.local
```

### 2. Check CoreDNS Logs:

Find CoreDNS pods:

```
kubectl get pods -n kube-system -l k8s-app=kube-dns
```

Examine logs for errors related to `project-beta` lookups:

```
kubectl logs -n kube-system <coredns-pod-name> -f
```

Look for errors like `NXDOMAIN`, `SERVFAIL`, or specific plugin errors.

### 3. Inspect Network Policies:

Check if any NetworkPolicy in `project-beta` restricts ingress from `project-alpha`:

```
kubectl get networkpolicy -n project-beta
```

Describe relevant policies to check `podSelector` and `ingress` rules:

```
kubectl describe networkpolicy <policy-name> -n project-beta
```



Check if any NetworkPolicy in `kube-system` restricts egress from CoreDNS pods to the API server or if egress from `project-alpha` to `kube-system` (UDP/TCP 53) is blocked.

```
kubectrl get networkpolicy -n kube-system
```

```
kubectrl get networkpolicy -n project-alpha
```

#### 4. Verify Service and Endpoint Existence in `project-beta`:

Ensure the target service exists and has endpoints:

```
kubectrl get svc beta-service -n project-beta
```

```
kubectrl get endpoints beta-service -n project-beta
```

If endpoints are missing, troubleshoot the pods backing `beta-service`.

#### 5. Review CoreDNS Configuration:

Examine the CoreDNS ConfigMap:

```
kubectrl get configmap coredns -n kube-system -o yaml
```

Look for incorrect `forward` directives, plugin configurations (`kubernetes`, `rewrite`, `acl`), or potential misconfigurations affecting cross-namespace lookups. Ensure the `kubernetes` plugin configuration is standard.

#### 6. Test DNS from a Different Namespace:

Exec into a pod in a third namespace (e.g., `default`) and attempt the lookup to rule out a `project-alpha`-specific issue:

```
kubectrl exec -it -n default <default-pod-name> -- nslookup beta-service.project-beta.svc.cluster.local
```

#### 7. Implement Fix (Example: Network Policy Adjustment):

If a NetworkPolicy in `project-beta` was blocking DNS lookups (CoreDNS needs to reach pods/endpoints), adjust it to allow ingress from `kube-system` namespace or specifically from CoreDNS pods:

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: allow-dns-queries
```



```
namespace: project-beta
spec:
  podSelector: {} # Apply to all pods in project-beta
  policyTypes:
  - Ingress
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          kubernetes.io/metadata.name: kube-system # Or a more specific label for kube-
system
      podSelector:
        matchLabels:
          k8s-app: kube-dns
    ports:
    - protocol: UDP
      port: 53
    - protocol: TCP # Less common for queries, but good practice
      port: 53
  # ... include other necessary ingress rules here ...
```

Apply the updated policy:

```
kubectl apply -f allow-dns-policy.yaml -n project-beta
```

## 8. Re-Verify DNS Resolution:

Repeat Step 1 from the `project-alpha` pod.

## Outcome:

Cross-namespace DNS resolution between `project-alpha` and `project-beta` is restored.



The root cause (e.g., overly restrictive NetworkPolicy, CoreDNS misconfiguration) is identified and rectified.

Inter-service communication dependent on DNS is functional again.

---

## **Scenario 127: StatefulSet PVC Binding Failure on Replica Scale-Up**

### **Scenario**

You scale up a StatefulSet named `event-processor` in the `project-stream` namespace from 3 to 5 replicas. The new pods (`event-processor-3` and `event-processor-4`) get stuck in the `Pending` state. Describing the pods reveals their PersistentVolumeClaims (PVCs) like `data-event-processor-3` cannot be bound. Existing replicas (`-0` to `-2`) are running fine with their PVCs bound.

### **Solution**

Context:

```
kubectl config use-context k8s-c18-stage
```

### **Steps:**

#### **1. Confirm Pod Status and Events:**

Check pod status:

```
kubectl get pods -n project-stream -l app=event-processor
```

Describe a pending pod to see events related to scheduling and volume binding:

```
kubectl describe pod event-processor-3 -n project-stream
```

Look for events like `FailedScheduling` (if node resources are the issue) or `FailedBinding` (specific to PVC). Note the PVC name (`data-event-processor-3`).

#### **2. Inspect the Pending PersistentVolumeClaim (PVC):**

Describe the specific PVC mentioned in the pod events:

```
kubectl describe pvc data-event-processor-3 -n project-stream
```

Look for events associated with the PVC itself. Common errors include:

```
`no persistent volumes available for this claim`
```



```
`storageclass.storage.k8s.io "<storage-class-name>" not found`
```

Errors from the external provisioner (if using dynamic provisioning).

### 3. Check Available Persistent Volumes (PVs):

List PVs and filter by the relevant StorageClass:

```
kubectl get pv --sort-by=.metadata.creationTimestamp
```

```
kubectl get pv -l storageclass=<storage-class-name> # Get StorageClass from PVC description
```

Look for PVs in `Available` state that match the PVC's requirements (StorageClass, access modes, capacity).

Check if existing PVs are already bound to the older PVCs (`data-event-processor-0` to `-2`).

### 4. Verify the StorageClass:

Ensure the StorageClass specified in the PVC (or the default one) exists and is correctly configured:

```
kubectl get storageclass <storage-class-name> -o yaml
```

Check the `provisioner` field, `reclaimPolicy`, `volumeBindingMode` (`WaitForFirstConsumer` can delay binding until pod scheduling), and any `parameters`.

### 5. Investigate the CSI Driver / External Provisioner:

If using dynamic provisioning, check the logs of the CSI controller/external provisioner pods (often in `kube-system` or a dedicated namespace):

```
kubectl get pods -n <csi-driver-namespace>
```

```
kubectl logs -n <csi-driver-namespace> <csi-provisioner-pod-name> -f
```

Look for errors related to volume creation, API credentials, quota limits, or storage backend issues.

### 6. Check Node-Specific Factors (if `volumeBindingMode: WaitForFirstConsumer`):

If the StorageClass uses `WaitForFirstConsumer`, binding waits until a pod is scheduled. Check if scheduling is failing due to node resource constraints, taints/tolerations, or affinity rules:

```
kubectl describe pod event-processor-3 -n project-stream # Look for scheduling failures
```

```
kubectl get nodes # Check node conditions and capacity
```

Also check CSI node plugin logs on potential target nodes for attachment errors.



## 7. Resolve the Issue (Example: Insufficient Quota in Provisioner):

Assume CSI logs show errors like "Quota Exceeded" on the storage backend.

Action: Increase the storage quota on the backend system (e.g., AWS EBS, GCE PD, vSphere datastore).

Kubernetes doesn't need direct changes, but the provisioner should now succeed on its next retry.

## 8. Verify PVC Binding and Pod Startup:

Monitor the PVC status:

```
kubectrl get pvc data-event-processor-3 -n project-stream -w
```

Once it changes to `Bound`, check the pod status:

```
kubectrl get pods event-processor-3 -n project-stream -w
```

The pod should transition from `Pending` to `ContainerCreating` and then `Running`. Repeat for `event-processor-4`.

## Outcome:

The root cause of the PVC binding failure (e.g., lack of available PVs, StorageClass issue, provisioner error, quota limit) is identified and resolved.

The new StatefulSet replicas (`event-processor-3`, `event-processor-4`) successfully bind their PVCs.

The pods start correctly, and the StatefulSet reaches the desired scale of 5 replicas.

---

## **Scenario 128: API Server Performance Degradation Under Load**

### Scenario

Users report intermittent slowness and timeouts when using `kubectrl`. CI/CD pipelines interacting with the cluster are failing sporadically. Monitoring reveals high latency and increased CPU/Memory usage for `kube-apiserver` pods, but certificates are valid, and etcd seems healthy (`etcdctl endpoint health` is OK). The cluster isn't throwing obvious errors, just becoming sluggish.



## Solution

Context:

```
kubectl config use-context k8s-c19-perf
```

Steps:

### 1. Confirm API Server Slowness:

Run simple, repeated `kubectl` commands and measure response time:

```
time kubectl get nodes
```

```
time kubectl get pods -A -l app=some-common-label --limit 10
```

Compare response times during peak vs. off-peak hours if possible.

### 2. Monitor API Server Resources and Metrics:

Check resource utilization of API server pods:

```
kubectl top pod -n kube-system -l component=kube-apiserver
```

Use Prometheus/Grafana (or similar) to examine key API server metrics:

`apiserver\_request\_duration\_seconds\_bucket`: High latency distribution?

`apiserver\_request\_total`: High overall request rate? Filter by `verb`, `resource`, `client`.

`apiserver\_current\_inflight\_requests`: High number of concurrent requests?

`workqueue\_depth`: For controllers, indicates potential processing delays.

`etcd\_request\_duration\_seconds\_bucket`: Rule out etcd latency being the primary cause (even if health is okay).

### 3. Identify Heavy API Clients/Controllers:

Analyze `apiserver\_request\_total` broken down by `client` (User Agent) and `resource`/`verb`.

Look for specific clients (e.g., a custom controller, monitoring agent, CI/CD system) making excessive `LIST` or `WATCH` calls, especially on large resource sets (Pods, Events, ConfigMaps).

Check API server audit logs for frequent requests from specific IPs or service accounts if enabled.





#### 4. Inspect API Server Logs:

Check logs for throttling messages or specific errors:

```
kubectl logs -n kube-system <kube-apiserver-pod-name> | grep -i  
'throttling\|timeout\|error'
```

#### 5. Review Resource Limits/Requests for API Server:

Check the manifest used to deploy the API server (often a static pod manifest in `/etc/kubernetes/manifests/kube-apiserver.yaml` on control plane nodes):

```
cat /etc/kubernetes/manifests/kube-apiserver.yaml | grep -A2 resources:
```

Are the CPU/memory requests and limits adequate for the observed load?

#### 6. Consider API Priority and Fairness (APF):

Check if APF is enabled (default in recent versions):

```
kubectl get flowschemas
```

```
kubectl get prioritylevelconfigurations
```

APF might be throttling less important requests to protect the server, which could be perceived as slowness by those clients. Check metrics related to APF (e.g., `apiserver\_flowcontrol\_request\_wait\_duration\_seconds`).

#### 7. Optimize Problematic Clients:

If a specific client/controller is identified as the source of excessive load:

Reduce LIST/WATCH scope: Use label selectors, field selectors, or watch specific objects instead of entire collections.

Increase sync intervals: Don't reconcile resources too frequently if not needed.

Use Informers/Caches: Ensure controllers use shared informers to cache objects locally instead of querying the API server repeatedly.

Implement client-side throttling: Limit the rate of requests the client makes.

#### 8. Scale Up/Tune API Server (If Necessary):

If the overall legitimate load is high, consider:

Increasing CPU/Memory limits and requests for the API server pods.



Scaling out the number of API server replicas (requires load balancing).

Tuning API server flags (e.g., `--max-requests-inflight`, `--max-mutating-requests-inflight`), use with caution.

## 9. Monitor Post-Changes:

Continuously monitor API server metrics (latency, resource usage) after applying fixes (client optimization or server scaling).

Verify `kubectl` responsiveness and CI/CD pipeline stability.

## Outcome

The bottleneck causing API server slowness is identified (e.g., a misbehaving controller, insufficient resources, high legitimate load).

Corrective actions are taken, such as optimizing API clients or scaling/tuning the API server.

API server performance and responsiveness are restored to acceptable levels.

Cluster interaction stability (`kubectl`, CI/CD) is improved.

---

## **Scenario 129: CronJob Failing to Execute Scheduled Tasks Reliably**

### Scenario

A critical nightly CronJob named `daily-report` in the `reporting` namespace, scheduled to run at `0 2 * * *` (2 AM daily), has become unreliable. Sometimes it runs, creating a Job pod, but other times it completely misses its schedule, and no Job is created. There are no obvious errors in the CronJob controller logs.

### Solution

Context:

```
kubectl config use-context k8s-c20-batch
```

Steps:



## 1. Verify CronJob Definition and Status:

Get the full CronJob definition:

```
kubectl get cronjob daily-report -n reporting -o yaml
```

Check:

`schedule`: Is it `0 2`? Verify syntax correctness.

`suspend`: Is it `false` or unset? If `true`, the CronJob is paused.

`concurrencyPolicy`: (`Allow`, `Forbid`, `Replace`) - Could `Forbid` prevent runs if a previous job is stuck?

`successfulJobsHistoryLimit` / `failedJobsHistoryLimit`: Are old jobs being cleaned up?

`startingDeadlineSeconds`: Is it set? If too short, a missed schedule due to controller lag might cause the job to be skipped entirely.

## 2. Check Recent Job History:

List Jobs created by this CronJob:

```
kubectl get jobs -n reporting -l cronjob=daily-report --sort-by=.metadata.creationTimestamp
```

Note the timestamps of the created Jobs. Are there gaps corresponding to the missed schedules?

Check the status of recent Jobs (Succeeded, Failed). A failed job might impact future runs depending on `concurrencyPolicy`.

## 3. Examine Kubernetes Events:

Events often provide clues about scheduling decisions or failures:

```
kubectl get events -n reporting --field-selector involvedObject.kind=CronJob,involvedObject.name=daily-report
```

```
kubectl get events -n reporting --field-selector involvedObject.kind=Job # Look for creation/deletion events
```

Look for events like `SuccessfulCreate`, `FailedCreate`, or potentially scheduling issues related to the Job's pod template.

## 4. Check Controller Manager Logs:

The CronJob controller runs within the `kube-controller-manager`. Check its logs on the master node(s):



```
# Find the pod
```

```
kubectl get pods -n kube-system -l component=kube-controller-manager
```

```
# Check logs, grep for the CronJob name or namespace
```

```
kubectl logs -n kube-system <kube-controller-manager-pod-name> | grep 'daily-report\|reporting'
```

Look for errors related to creating Jobs, time synchronization issues, or controller loops being slow.

## 5. Investigate Resource Quotas:

Does the `reporting` namespace have ResourceQuotas that might prevent Job or Pod creation?

```
kubectl get resourcequota -n reporting
```

```
kubectl describe resourcequota <quota-name> -n reporting
```

Check if limits for `jobs.batch`, `pods`, `cpu`, or `memory` are being hit around 2 AM.

## 6. Verify RBAC Permissions for the Job's ServiceAccount:

Identify the ServiceAccount used by the Job template within the CronJob spec (`spec.jobTemplate.spec.template.spec.serviceAccountName`).

Ensure this ServiceAccount has the necessary permissions (e.g., create Pods, access Secrets/ConfigMaps) required by the Job's container. While this usually causes Job failure rather than non-creation, it's worth checking if related events are confusing.

```
kubectl describe rolebinding,clusterrolebinding -n reporting | grep <serviceAccountName>
```

```
# Or use a tool like rbac-lookup
```

## 7. Consider Time Synchronization:

Ensure all nodes (especially control plane nodes running `kube-controller-manager`) have accurate and synchronized time (e.g., using NTP). Significant time skew can disrupt schedules.

```
Check time on master node: `date`
```

## 8. Address the Likely Cause (Example: `startingDeadlineSeconds`):

Assume investigation suggests the controller manager was occasionally busy or delayed around 2 AM, and the default or a very short `startingDeadlineSeconds` caused the CronJob controller to skip the run if it couldn't create the Job very close to the scheduled time.



Solution: Add or increase `startingDeadlineSeconds` in the CronJob spec to allow more flexibility:

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: daily-report
  namespace: reporting
spec:
  schedule: "0 2 *"
  startingDeadlineSeconds: 120 # Allow 2 minutes past schedule before skipping
  jobTemplate:
    # ... rest of template
```

Apply the change:

```
kubectl apply -f daily-report-cronjob.yaml
```

## 9. Monitor Over Several Cycles:

Observe the CronJob over the next few days, checking Job history (Step 2) to confirm it now runs reliably at 2 AM.

### Outcome:

The reason for the CronJob's unreliable execution (e.g., missed schedule deadline, resource quota exhaustion, controller lag) is identified.

The CronJob configuration is adjusted (e.g., adding `startingDeadlineSeconds`, clearing stuck jobs, increasing quotas).

The `daily-report` CronJob now reliably creates its associated Job at the scheduled time.

-----



## **Scenario 130: Ingress Routing Failure for a Specific Path or Service**

### **Scenario**

An Ingress resource is configured to route traffic to multiple backend services. Requests to `https://app.example.com/serviceA` work correctly, mapping to `service-a` in the `apps` namespace. However, requests to `https://app.example.com/serviceB` consistently return a 503 Service Unavailable or 404 Not Found error from the Ingress controller, even though `service-b` and its pods in the `apps` namespace appear healthy.

### **Solution**

Context:

```
kubectrl config use-context k8s-c21-gateway
```

Steps:

#### **1. Verify Pod and Service Health for `service-b`:**

Check pods backing `service-b`:

```
kubectrl get pods -n apps -l app=service-b -o wide
```

Ensure they are `Running` and `Ready`. Check their logs for internal errors.

Check the Service definition:

```
kubectrl get svc service-b -n apps -o yaml
```

Verify `selector`, `port`, and `targetPort`.

Check the Endpoints object associated with the service:

```
kubectrl get endpoints service-b -n apps
```

Ensure it lists the correct IP addresses and ports of the healthy `service-b` pods. If empty or incorrect, troubleshoot the pod selectors or pod readiness.

#### **2. Inspect the Ingress Resource Definition:**

Get the full Ingress definition:

```
kubectrl get ingress main-ingress -n apps -o yaml
```

Carefully review the `rules` and `paths` sections:

Is there a rule for `host: app.example.com`?



Under that host, is there a `path` entry for `/serviceB`?

Is the `pathType` correct (`Prefix`, `Exact`, `ImplementationSpecific`)? Mismatches are common causes. (`Prefix` is often safest).

Does the `backend` for `/serviceB` correctly specify `service.name: service-b` and `service.port.name` (or `number`) matching the Service definition? Check for typos.

Are there any relevant annotations specific to your Ingress controller (e.g., rewrite rules, backend protocol) that might be misconfigured for `/serviceB`?

### 3. Check Ingress Controller Logs:

Find the Ingress controller pods (e.g., `nginx-ingress`, `traefik`, `haproxy-ingress` - often in a dedicated namespace like `ingress-nginx` or `kube-system`):

```
kubectl get pods -n <ingress-controller-namespace>
```

Tail the logs while sending a request to `https://app.example.com/serviceB`:

```
kubectl logs -n <ingress-controller-namespace> <ingress-controller-pod-name> -f
```

Look for log entries corresponding to the failed request. They often indicate:

Which upstream (service endpoint) it tried to connect to.

Connection errors (`connect() failed`, `timeout`).

Configuration errors (`could not find backend`, `rule syntax error`).

HTTP status codes returned from the backend service (if the connection succeeded but the app failed).

### 4. Verify Network Connectivity from Ingress Controller to Pods:

If logs suggest connection issues:

Check Network Policies in the `apps` namespace that might block ingress from the Ingress controller's namespace/pods.

```
kubectl get networkpolicy -n apps
```

Exec into the Ingress controller pod and try to reach `service-b`'s ClusterIP or pod IPs directly (using `curl`, `wget`):

```
kubectl exec -it -n <ingress-controller-namespace> <ingress-controller-pod-name> -- /bin/sh
```

```
curl -v http://<service-b-cluster-ip>:<service-port>/
```

```
curl -v http://<service-b-pod-ip>:<target-port>/ # Try a known pod IP
```



## 5. Examine Ingress Controller Configuration:

Some controllers generate internal configuration (e.g., `nginx.conf` for nginx-ingress). If possible, inspect this generated config inside the controller pod to see how it interpreted the Ingress resource.

```
# Example for nginx-ingress
```

```
kubectl exec -it -n <ingress-ns> <nginx-pod> -- cat /etc/nginx/nginx.conf
```

Look for the server block corresponding to `app.example.com` and the location block for `/serviceB`.

## 6. Address the Likely Cause (Example: Incorrect `pathType`):

Assume the Ingress definition used `pathType: Exact` for `/serviceB`, but requests were coming in as `/serviceB/` (with a trailing slash). `Exact` matching would fail.

Solution: Change `pathType` to `Prefix` in the Ingress definition:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: main-ingress
```

```
  namespace: apps
```

```
spec:
```

```
  rules:
```

```
    - host: app.example.com
```

```
      http:
```

```
        paths:
```

```
          - path: /serviceA
```

```
            pathType: Prefix # Assuming this was already Prefix
```

```
          backend:
```

```
            service:
```

```
              name: service-a
```

```
            port:
```

```
              number: 80
```

```
          - path: /serviceB # Path that was failing
```





```
pathType: Prefix # Changed from Exact
backend:
  service:
    name: service-b
    port:
      number: 8080
# ... tls section ...
```

Apply the change:

```
kubectl apply -f main-ingress.yaml
```

## 7. Test the Endpoint Again:

Send a request to `https://app.example.com/serviceB` (and maybe `https://app.example.com/serviceB/`):

```
curl -k https://app.example.com/serviceB # Use -k if using self-signed certs
```

Expect a successful response (e.g., 200 OK) from `service-b`.

## Outcome:

The specific misconfiguration preventing Ingress routing to `service-b` (e.g., incorrect `path`, `pathType`, service name/port, NetworkPolicy block, backend pod issue) is found and corrected.

Requests to `https://app.example.com/serviceB` are now successfully routed by the Ingress controller to the healthy `service-b` pods.

Full application functionality dependent on both `/serviceA` and `/serviceB` paths is restored.

Stay tuned for Part 27, where we will continue navigating the complexities of Kubernetes troubleshooting! Follow for more insights.



In the up-coming parts, we will discussion on more troubleshooting steps for the different Kubernetes based scenarios. So, stay tuned for the and follow @Prasad Suman Mohan for more such posts.

